

- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. 2018. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv:1803.05457v1*.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. [Training Verifiers to Solve Math Word Problems](#). *arxiv:2110.14168[cs]*.
- Owen Colegrove, Vik Paruchuri, and OpenPhi-Team. 2024. [OpenPhi/textbooks · Datasets at Hugging Face](#).
- Róbert Csordás, Kazuki Irie, Jürgen Schmidhuber, Christopher Potts, and Christopher D. Manning. 2024. [MoEUT: Mixture-of-Experts Universal Transformers](#). In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
- Gautier Dagan. 2024. [Bpeasy](#).
- Gautier Dagan, Gabriel Synnaeve, and Baptiste Rozière. 2024. [Getting the most out of your tokenizer for pre-training and domain adaptation](#). *arxiv:2402.01035[cs]*.
- Tri Dao. 2023. [FlashAttention-2: Faster Attention with Better Parallelism and Work Partitioning](#). *arxiv:2307.08691[cs]*.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. 2022. [FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness](#). *arxiv:2205.14135[cs]*.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyi Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, and 181 others. 2025. [DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning](#). *arxiv:2501.12948[cs]*.
- DeepSeek-AI, Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Daya Guo, Dejian Yang, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, and 181 others. 2024. [DeepSeek-V3 Technical Report](#). *arxiv:2412.19437[cs]*.
- Mostafa Dehghani, Stephan Gouws, Oriol Vinyals, Jakob Uszkoreit, and Łukasz Kaiser. 2019. [Universal Transformers](#). *arxiv:1807.03819[cs, stat]*.
- Yuntian Deng, Yejin Choi, and Stuart Shieber. 2024. [From Explicit CoT to Implicit CoT: Learning to Internalize CoT Step by Step](#). *arxiv:2405.14838[cs]*.
- Hantian Ding, Zijian Wang, Giovanni Paolini, Varun Kumar, Anoop Deoras, Dan Roth, and Stefano Soatto. 2024. [Fewer Truncations Improve Language Modeling](#). In *Forty-First International Conference on Machine Learning*.
- Ming Ding, Zhuoyi Yang, Wenyi Hong, Wendi Zheng, Chang Zhou, Da Yin, Junyang Lin, Xu Zou, Zhou Shao, Hongxia Yang, and Jie Tang. 2021. [CogView: Mastering Text-to-Image Generation via Transformers](#). In *Advances in Neural Information Processing Systems*, volume 34, pages 19822–19835. Curran Associates, Inc.
- Maha Elbayad, Jiatao Gu, Edouard Grave, and Michael Auli. 2019. [Depth-Adaptive Transformer](#). In *International Conference on Learning Representations*.
- Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, Ahmed A. Aly, Beidi Chen, and Carole-Jean Wu. 2024. [LayerSkip: Enabling Early Exit Inference and Self-Speculative Decoding](#). *arxiv:2404.16710[cs]*.
- Katie Everett, Lechao Xiao, Mitchell Wortsman, Alexander A. Alemi, Roman Novak, Peter J. Liu, Izzeddin Gur, Jascha Sohl-Dickstein, Leslie Pack Kaelbling, Jaehoon Lee, and Jeffrey Pennington. 2024. [Scaling Exponents Across Parameterizations and Optimizers](#). *arxiv:2407.05872[cs]*.
- Angela Fan, Edouard Grave, and Armand Joulin. 2019. [Reducing Transformer Depth on Demand with Structured Dropout](#). *arxiv:1909.11556[cs, stat]*.
- Angela Fan, Thibaut Lavril, Edouard Grave, Armand Joulin, and Sainbayar Sukhbaatar. 2021. [Addressing Some Limitations of Transformers with Feedback Memory](#). *arxiv:2002.09402[cs, stat]*.
- Ying Fan, Yilun Du, Kannan Ramchandran, and Kangwook Lee. 2025. [Looped Transformers for Length Generalization](#). In *The Thirtieth International Conference on Learning Representations*.
- William Fedus, Barret Zoph, and Noam Shazeer. 2022. [Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity](#). *arxiv:2101.03961[cs]*.
- Xidong Feng, Yicheng Luo, Ziyang Wang, Hongrui Tang, Mengyue Yang, Kun Shao, David Mguni, Yali Du, and Jun Wang. 2023. [ChessGPT: Bridging Policy Learning and Language Modeling](#). *Advances in Neural Information Processing Systems*, 36:7216–7262.
- Sebastian Gabarain. 2024. [Locutusque/hercules-v5.0 · Datasets at Hugging Face](#).
- Khashayar Gatmiry, Nikunj Saunshi, Sashank J. Reddi, Stefanie Jegelka, and Sanjiv Kumar. 2024. [Can Looped Transformers Learn to Implement Multi-step Gradient Descent for In-context Learning?](#)
- Jonas Geiping and Tom Goldstein. 2023. [Cramming: Training a Language Model on a single GPU in one day](#). In *Proceedings of the 40th International Conference on Machine Learning*, pages 11117–11143. PMLR.
- Jonas Geiping and Michael Moeller. 2019. [Parametric Majorization for Data-Driven Energy Minimization Methods](#). In *Proceedings of the IEEE International Conference on Computer Vision*, pages 10262–10273.
- F.A. Gers and J. Schmidhuber. 2000. [Recurrent nets that time and count](#). In *Proceedings of the IEEE-INNS-ENNS International Joint Conference on Neural Networks. IJCNN 2000. Neural Computing: New Challenges and Perspectives for the New Millennium*, volume 3, pages 189–194 vol.3.
- Angeliki Giannou, Shashank Rajput, Jy-Yong Sohn, Kangwook Lee, Jason D. Lee, and Dimitris Papailiopoulos. 2023. [Looped Transformers as Programmable Computers](#). In *Proceedings of the 40th International Conference on Machine Learning*, pages 11398–11442. PMLR.

- Priya Goyal, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. 2018. [Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour](#). *arxiv:1706.02677[cs]*.
- Alex Graves. 2017. [Adaptive Computation Time for Recurrent Neural Networks](#). *arxiv:1603.08983[cs]*.
- Alex Graves, Greg Wayne, and Ivo Danihelka. 2014. [Neural Turing Machines](#). *arxiv:1410.5401[cs]*.
- Dirk Groeneveld, Iz Beltagy, Pete Walsh, Akshita Bhagia, Rodney Kinney, Oyvind Tafjord, Ananya Harsh Jha, Hamish Ivison, Ian Magnusson, Yizhong Wang, Shane Arora, David Atkinson, Russell Authur, Khyathi Raghavi Chandu, Arman Cohan, Jennifer Dumas, Yanai Elazar, Yuling Gu, Jack Hessel, and 24 others. 2024. [OLMo: Accelerating the Science of Language Models](#). *arxiv:2402.00838[cs]*.
- Alexander Hägele, Elie Bakouch, Atli Kosson, Loubna Ben Allal, Leandro Von Werra, and Martin Jaggi. 2024. [Scaling Laws and Compute-Optimal Training Beyond Fixed Training Durations](#). In *Workshop on Efficient Systems for Foundation Models II @ ICML2024*.
- Shibo Hao, Sainbayar Sukhbaatar, DiJia Su, Xian Li, Zhiting Hu, Jason Weston, and Yuandong Tian. 2024. [Training Large Language Models to Reason in a Continuous Latent Space](#). *arxiv:2412.06769[cs]*.
- Tamir David Hay and Lior Wolf. 2023. [Dynamic Layer Tying for Parameter-Efficient Transformers](#). In *The Twelfth International Conference on Learning Representations*.
- Zexue He, Leonid Karlinsky, Donghyun Kim, Julian McAuley, Dmitry Krotov, and Rogerio Feris. 2024. [CAMELoT: Towards Large Language Models with Training-Free Consolidated Associative Memory](#). *arxiv:2402.13449[cs]*.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021a. Measuring massive multitask language understanding. *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021b. [Measuring Massive Multitask Language Understanding](#). In *International Conference on Learning Representations*.
- J J Hopfield. 1982. [Neural networks and physical systems with emergent collective computational abilities](#). *Proceedings of the National Academy of Sciences of the United States of America*, 79(8):2554–2558.
- Jiewen Hu, Thomas Zhu, and Sean Welleck. 2024. [miniCTX: Neural Theorem Proving with \(Long-\)Contexts](#). *arxiv:2408.03350[cs]*.
- Pavel Izmailov, Dmitrii Podoprikin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson. 2018. [Averaging weights leads to wider optima and better generalization: 34th Conference on Uncertainty in Artificial Intelligence 2018, UAI 2018, 34th Conference on Uncertainty in Artificial Intelligence 2018, UAI 2018](#), pages 876–885.
- Albert Q. Jiang, Wenda Li, and Mateja Jamnik. 2023. [Multilingual Mathematical Autoformalization](#). *arxiv:2311.03755[cs]*.
- Matt Gardner Johannes Welbl, Nelson F. Liu. 2017. Crowdsourcing multiple choice science questions.
- Jean Kaddour. 2022. [Stop Wasting My Time! Saving Days of ImageNet and BERT Training with Latest Weight Averaging](#). *arxiv:2209.14981[cs, stat]*.
- Guy Kaplan, Matanel Oren, Yuval Reif, and Roy Schwartz. 2024. [From Tokens to Words: On the Inner Lexicon of LLMs](#). *arxiv:2410.05864[cs]*.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. [Scaling Laws for Neural Language Models](#). *arxiv:2001.08361[cs, stat]*.
- Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. 2020. [Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention](#). In *Proceedings of the 37th International Conference on Machine Learning*, pages 5156–5165. PMLR.
- Matthew Kenney. 2024. [ArXivDLInstruct](#).
- Seungone Kim, Juyoung Suk, Shayne Longpre, Bill Yuchen Lin, Jamin Shin, Sean Welleck, Graham Neubig, Moontae Lee, Kyungjae Lee, and Minjoon Seo. 2024. [Prometheus 2: An Open Source Language Model Specialized in Evaluating Other Language Models](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4334–4353, Miami, Florida, USA. Association for Computational Linguistics.
- Diederik P. Kingma and Jimmy Ba. 2015. [Adam: A Method for Stochastic Optimization](#). In *International Conference on Learning Representations (ICLR)*, San Diego.
- Wojciech Kryściński, Nazneen Rajani, Divyansh Agarwal, Caiming Xiong, and Dragomir Radev. 2022. [BookSum: A Collection of Datasets for Long-form Narrative Summarization](#). *arxiv:2105.08209[cs]*.
- Xin Lai, Zhuotao Tian, Yukang Chen, Senqiao Yang, Xian-gru Peng, and Jiaya Jia. 2024. [Step-DPO: Step-wise Preference Optimization for Long-chain Reasoning of LLMs](#). *arxiv:2406.18629[cs]*.
- Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2019. [ALBERT: A Lite BERT for Self-supervised Learning of Language Representations](#). In *International Conference on Learning Representations*.
- Yann LeCun. 2022. A Path Towards Autonomous Machine Intelligence. *Preprint*, Version 0.9.2:62.
- Yann LeCun and Fu Jie Huang. 2005. [Loss functions for discriminative training of energy-based models](#). In *AISTATS 2005 - Proceedings of the 10th International Workshop on Artificial Intelligence and Statistics*, pages 206–213.
- Jason Lee, Elman Mansimov, and Kyunghyun Cho. 2018. [Deterministic Non-Autoregressive Neural Sequence Modeling by Iterative Refinement](#). In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1173–1182, Brussels, Belgium. Association for Computational Linguistics.

- Yaniv Leviathan, Matan Kalman, and Yossi Matias. 2023. [Fast Inference from Transformers via Speculative Decoding](#). In *Proceedings of the 40th International Conference on Machine Learning*, pages 19274–19286. PMLR.
- Yoav Levine, Noam Wies, Or Sharir, Hofit Bata, and Amnon Shashua. 2021. [The Depth-to-Width Interplay in Self-Attention](#). *arxiv:2006.12467[cs, stat]*.
- Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. 2022. [Solving quantitative reasoning problems with language models](#). *Preprint*, arXiv:2206.14858.
- Raymond Li, Loubna Ben Allal, Yangtian Zi, Niklas Muennighoff, Denis Kocetkov, Chenghao Mou, Marc Marone, Christopher Akiki, Jia Li, Jenny Chim, Qian Liu, Evgenii Zheltonozhskii, Terry Yue Zhuo, Thomas Wang, Olivier Dehaene, Joel Lamy-Poirier, Joao Monteiro, Nicolas Gontier, Ming-Ho Yee, and 39 others. 2023. [StarCoder: May the source be with you!](#) *Transactions on Machine Learning Research*.
- Xian Li, Asa Cooper Stickland, Yuqing Tang, and Xiang Kong. 2020a. [Deep Transformers with Latent Depth](#). *arxiv:2009.13102[cs]*.
- Yujia Li, Felix Gimeno, Pushmeet Kohli, and Oriol Vinyals. 2020b. [Strong Generalization and Efficiency in Neural Programs](#). *arxiv:2007.03629[cs]*.
- Omkar Pangarkar Liping Tang, Nikhil Ranjan. 2024. [TxT360: A top-quality LLM pre-training dataset requires the perfect blend](#).
- Hao Liu, Matei Zaharia, and Pieter Abbeel. 2023a. [Ring attention with blockwise transformers for near-infinite context](#). *arXiv preprint arXiv:2310.01889*.
- Luyang Liu, Jonas Pfeiffer, Jiaxing Wu, Jun Xie, and Arthur Szlam. 2024a. [Deliberation in Latent Space via Differentiable Cache Augmentation](#). *arxiv:2412.17747[cs]*.
- Xiao Liu, Hanyu Lai, Hao Yu, Yifan Xu, Aohan Zeng, Zhengxiao Du, Peng Zhang, Yuxiao Dong, and Jie Tang. 2023b. [WebGLM: Towards An Efficient Web-Enhanced Question Answering System with Human Preferences](#). In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD '23*, pages 4549–4560, New York, NY, USA. Association for Computing Machinery.
- Zechun Liu, Changsheng Zhao, Forrest Iandola, Chen Lai, Yuandong Tian, Igor Fedorov, Yunyang Xiong, Ernie Chang, Yangyang Shi, Raghuraman Krishnamoorthi, Liangzhen Lai, and Vikas Chandra. 2024b. [MobileLLM: Optimizing Sub-billion Parameter Language Models for On-Device Use Cases](#). *arxiv:2402.14905[cs]*.
- Zhengzhong Liu, Aurick Qiao, Willie Neiswanger, Hongyi Wang, Bowen Tan, Tianhua Tao, Junbo Li, Yuqi Wang, Suqi Sun, Omkar Pangarkar, Richard Fan, Yi Gu, Victor Miller, Yonghao Zhuang, Guowei He, Haonan Li, Fajri Koto, Liping Tang, Nikhil Ranjan, and 9 others. 2023c. [LLM360: Towards fully transparent open-source LLMs](#).
- Ilya Loshchilov and Frank Hutter. 2017. [Decoupled Weight Decay Regularization](#). *arXiv:1711.05101 [cs, math]*.
- Anton Lozhkov, Raymond Li, Loubna Ben Allal, Federico Cas-sano, Joel Lamy-Poirier, Nouamane Tazi, Ao Tang, Dmytro Pykhtar, Jiawei Liu, Yuxiang Wei, Tianyang Liu, Max Tian, Denis Kocetkov, Arthur Zucker, Younes Belkada, Zijian Wang, Qian Liu, Dmitry Abulkhanov, Indraneil Paul, and 47 others. 2024. [StarCoder 2 and The Stack v2: The Next Generation](#).
- Zimu Lu, Aojun Zhou, Ke Wang, Houxing Ren, Weikang Shi, Junting Pan, Mingjie Zhan, and Hongsheng Li. 2024. [MathCoder2: Better Math Reasoning from Continued Pretraining on Model-translated Mathematical Code](#). *arxiv:2410.08196[cs]*.
- Sebastian Majstorovic. 2024. [Selected Digitized Books | The Library of Congress](#).
- Larisa Markeeva, Sean McLeish, Borja Ibarz, Wilfried Bounsi, Olga Kozlova, Alex Vitvitskyi, Charles Blundell, Tom Goldstein, Avi Schwarzschild, and Petar Veličković. 2024. [The CLRS-Text Algorithmic Reasoning Language Benchmark](#). *arxiv:2406.04229[cs]*.
- Mrinal Mathur, Barak A. Pearlmutter, and Sergey M. Plis. 2024. [MIND over Body: Adaptive Thinking using Dynamic Computation](#). In *The Thirteenth International Conference on Learning Representations*.
- Sean Michael McLeish, Arpit Bansal, Alex Stein, Neel Jain, John Kirchenbauer, Brian R. Bartoldson, Bhavya Kailkhura, Abhinav Bhatele, Jonas Geiping, Avi Schwarzschild, and Tom Goldstein. 2024. [Transformers Can Do Arithmetic with the Right Embeddings](#). In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.
- William Merrill, Ashish Sabharwal, and Noah A. Smith. 2022. [Saturated Transformers are Constant-Depth Threshold Circuits](#). *Transactions of the Association for Computational Linguistics*, 10:843–856.
- Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. 2018. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *EMNLP*.
- Tomáš Mikolov, Stefan Kombrink, Lukáš Burget, Jan Černocký, and Sanjeev Khudanpur. 2011. [Extensions of recurrent neural network language model](#). In *2011 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 5528–5531.
- Takeru Miyato, Sindy Löwe, Andreas Geiger, and Max Welling. 2024. [Artificial Kuramoto Oscillatory Neurons](#). In *The Thirteenth International Conference on Learning Representations*, Singapore.
- Ryan Moulton. 2023. [The Many Ways that Digital Minds Can Know](#).
- Niklas Muennighoff, Qian Liu, Armel Zebaze, Qinkai Zheng, Binyuan Hui, Terry Yue Zhuo, Swayam Singh, Xiangru Tang, Leandro von Werra, and Shayne Longpre. 2024. [OctoPack: Instruction Tuning Code Large Language Models](#). *arxiv:2308.07124[cs]*.
- Nam Pham. 2023. [Tiny-textbooks \(Revision 14de7ba\)](#).
- Nam Pham. 2024. [Tiny-strange-textbooks \(Revision 6f304f1\)](#).